

COMPREHENSIVE ARCHITECTURE DOCUMENT

AI Chatbot, Recommendation Engine & Kundli PDF System

Complete end-to-end technical documentation covering Groq LLM orchestration, Gorse ML collaborative filtering, Astrological Soft-Boosting, Explicit Reason Badges, Kundli-to-Recommendation pipeline, 18+ D-Chart Kundli PDF Engine, telemetry event ingestion, system resilience, and all API specifications.

DOCUMENT VERSION

v7.0.0 (Master Edition)

LAST UPDATED

July 28, 2026

SYSTEM READINESS

100% Production Ready

Table of Contents


- 1. Executive Summary & Architecture Overview — Page 3
- 2. AI Chatbot (AstroGuide) Architecture — Page 4
- 3. Chatbot Capabilities & Feature Breakdown — Page 5
- 4. Chatbot System Prompt Engineering — Page 6
- 5. Chatbot Consultation Handoff & Guardrails — Page 7
- 6. Hybrid Recommendation Engine Architecture — Page 8
- 7. Gorse ML Engine & Collaborative Filtering — Page 9
- 8. Cold Start 3-Tier Fallback Strategy — Page 10
- 9. Astrological Soft-Boosting & Kundli Scoring — Page 11
- 10. Mahadasha-to-Category Mapping (Full Reference) — Page 12
- 11. Explicit Reasoning Badges System (Point 9) — Page 13
- 12. Recommendation Surface Points (Shop, PDP, Chat) — Page 14
- 13. Kundli PDF Generation Engine (18+ D-Charts) — Page 15
- 14. Kundli Astronomical Calculations & Avakhada Chakra — Page 16
- 15. Telemetry & Event Ingestion Pipeline — Page 17
- 16. Complete API Reference — Page 18
- 17. Integration Points & Tech Stack — Page 19
- 18. System Boundaries, Edge Cases & Resilience — Page 20
- 19. Final System Readiness Sign-Off — Page 20

1. Executive Summary & Architecture Overview

The Aroham platform operates as a hybrid microservices glue layer connecting a React 19 web application, an Express.js Node.js backend, the Groq LLM API, the Gorse ML Engine (on Docker), and the FreeAstroAPI / VedAstro astronomical calculation engine.

This Option B architecture replaces complex in-house model training with high-speed managed platform APIs, providing enterprise-grade scalability while maintaining deep Vedic customization through custom Node.js microservices.

1.1 High-Level System Topology



1.2 Backend Route Map

All backend routes are registered in `server.js` under the Express application. The server listens on port 5000 with full CORS support, 50MB request body limit (for image payloads), raw body preservation (for Razorpay webhook signature verification), and security headers (`X-Content-Type-Options` , `X-Frame-Options` , `Cache-Control`).

ROUTE	MODULE	PURPOSE
/api/chat	routes/chat.js	AI Chatbot (Groq LLM)
/api/recommendations	routes/recommendations.js	ML Recommendation Engine
/api/kundli	routes/kundli.js	Kundli PDF Generation
/api/telemetry	routes/telemetry.js	Behavioral Event Ingestion
/api/products	routes/products.js	Product Catalog
/api/cart	routes/cart.js	Shopping Cart

ROUTE	MODULE	PURPOSE
/api/orders	routes/orders.js	Order Management
/api/payments	routes/payments.js	Razorpay Integration
/api/addresses	routes/addresses.js	User Addresses
/api/auth	routes/auth.js	Authentication
/api/shiprocket	routes/shiprocket.js	Shipping & Logistics
/api/health	inline	Health Check Endpoint

2. AI Chatbot (AstroGuide) Architecture

AstroGuide is an empathetic, spiritually aligned AI chat companion powered by Groq's LLaMA 3.3 70B model. It handles user queries, integrates real-time product recommendations, and enforces strict Vedic guidance guardrails.

2.1 Technical Implementation (`routes/chat.js`)

The chatbot is implemented as a single `POST /api/chat` endpoint that orchestrates multiple subsystems in sequence:

Step 1: Request Payload Parsing

Accepts `{ message, userId, pageContext, history }`. The `userId` defaults to `user_demo_123` if not provided. The `pageContext` captures the current URL path from `window.location.pathname`. The `history` array contains all previous conversation messages for multi-turn memory.

Step 2: Gorse ML Product Retrieval

Before generating any LLM response, the backend queries the Gorse ML engine at `GET /api/recommend/{userId}?n=3` to fetch the top 3 personalized product candidates for the active user. These product IDs are then hydrated from Supabase PostgreSQL with full product details (name, image, price, description).

Step 3: Google Drive Image URL Formatting

Product images stored as Google Drive share links are automatically converted to direct-access URLs using the `formatImageUrl()` helper. This extracts the Drive file ID via regex and formats it as `https://lh3.googleusercontent.com/d/{fileId}` for reliable rendering.

Step 4: Price Formatting

All prices are stored as integer paise in the database. The chatbot formats them for display as `₹{(price / 100).toFixed(2)}` before injecting into the system prompt and returning to the frontend.

2.2 Request/Response Execution Flow

1. **User opens chat** via floating action widget on any page.
2. **Frontend retrieves** `aroham_astro_chat_history` from `sessionStorage`.
3. **User submits a message**.
4. **Frontend fires** `POST /api/chat` with `{ message, history, pageContext, userId }`.
5. **Backend queries Gorse ML** for top 3 personalized product candidates.
6. **Backend queries Supabase** for full product details of recommended items.
7. **Backend checks global Kundali profile cache** for astrological context.
8. **Backend constructs system prompt** with guardrails, page context, product data, and astro profile.
9. **Backend invokes Groq LLM API** (`llama-3.3-70b-versatile`) with temperature 0.7.
10. **Backend scans for consultation keywords** and sets `consultation_handoff: true`.

3. Chatbot Capabilities & Feature Breakdown

The AstroChatWidget (`AstroChatWidget.tsx`) implements a comprehensive floating chat interface with the following features:

3.1 Multi-Turn Conversation Memory

Retains full conversation history across page navigations using `sessionStorage` under the key `aroham_astro_chat_history` . On initial load, the widget checks for existing history and restores it. All messages are serialized as JSON arrays with structure: `{ id, sender, text, products, consultationHandoff }` .

3.2 Page-Aware Contextual Responses

Dynamically reads `window.location.pathname` and sends it as `pageContext` to the backend. The system prompt includes this context so the LLM can tailor responses when users ask about a specific product they are currently viewing.

3.3 ML + Astrological Product Cards

Embeds top 3 personalized remedies as interactive mini-cards with product thumbnails, names, prices, "View Product" navigation button, and "Add to Cart" direct-purchase button. Each card uses the same Google Drive image formatting pipeline for consistent thumbnail rendering.

3.4 Add All to Cart Bundle

When 2+ products are recommended, a single-click "Add All Recommended to Cart (Save 10%)" button appears below the product cards. This iterates through all recommended products and calls `addToCart()` for each, integrating with the shared cart state management system.

3.5 Astrologer Handoff Card

When the AI detects consultation-worthy topics (Kundali, birth chart, astrologer, horoscope, dasha, matchmaking), it automatically renders a "Book Live Video Consultation" card. This card features a gradient UI with a Video icon, description text, and a CTA button that navigates to `/consult` .

3.6 Proactive Engagement Bubble

After 5 seconds of page load (if the user has not interacted), a floating tooltip bubble appears above the chat icon with the message: "Looking for personalized Vedic remedies or astrological guidance?" This tooltip includes a close button and a pointer arrow. A pulsing red notification dot appears on the chat icon.

3.7 Quick Suggestion Pills

For new conversations (2 or fewer messages), four one-click suggestion pills are displayed:

- **Career & Business** — Pre-fills "Recommend a remedy for Career & Business"
- **Love & Harmony** — Pre-fills "Recommend a remedy for Love & Harmony"
- **Rahu/Ketu Protection** — Pre-fills "What is a good protection from Rahu/Ketu Protection"
- **Puja & Yantras** — Pre-fills "Tell me about Puja & Yantras"

3.8 Conversation Reset & Copy

"Start Fresh" button (RotateCcw icon) clears all messages, resets to the welcome message, and removes `sessionStorage` data. Each bot message includes a hover-reveal "Copy" button for one-click clipboard copy.

3.9 Guest User Persistence

For non-authenticated users, the widget generates a persistent `guest_{random8chars}` identifier stored in `localStorage` under `aroham_guest_user_id`. This ensures Gorse ML telemetry tracking works across sessions even without login.

4. Chatbot System Prompt Engineering

The system prompt is constructed dynamically in `routes/chat.js` with up to 6 contextual injection layers:

4.1 Base Identity Prompt

Role: "You are Aroham's Sacred AI AstroGuide — an empathetic, deeply knowledgeable Vedic astrological advisor."

The base prompt establishes the AI's identity, tone, and behavioral constraints across 4 key instructions:

- **Greeting & Tone:** Begin with "Namaste" or "Hari Om". Speak with warmth, wisdom, and respect.
- **No Guarantees or Doom:** Never make deterministic claims. Reframe remedies as supportive spiritual tools.
- **Sacred Terminology:** Use "temple-energized", "Vedic remedies", and "blessed tools for spiritual harmony".
- **Structuring:** Keep paragraphs short, clear, and easy to read. Use bullet points where appropriate.

4.2 Page Context Injection (Layer 5)

When the user is on a specific product page, the system prompt includes: "The user is currently viewing the page at {pathname}. If this is a specific product page, factor this context into your response as they might be asking about it."

4.3 Product Recommendation Injection

When Gorse ML returns product candidates, the system prompt appends a structured product list with names, prices, and descriptions. The LLM is instructed: "You MUST naturally weave a recommendation for 1 or 2 of these items into your guidance. Dedicate a section titled 'Recommended Sacred Remedies' at the bottom."

4.4 Astrological Profile Injection (Layer 6)

If the user has generated a Kundli PDF, their derived astrological profile is injected: "This devotee's active Mahadasha is {mahadasha} (Moon Sign: {moonSign}, Nakshatra: {nakshatra}). {guidanceNote}." This gives the LLM deep personalization context.

4.5 LLM API Configuration

PARAMETER	VALUE	RATIONALE
Provider	Groq Cloud API	Lowest latency LLM inference (sub-second)
Model	llama-3.3-70b-versatile	Best quality for nuanced spiritual guidance
Temperature	0.7	Balance between creativity and consistency

PARAMETER	VALUE	RATIONALE
Endpoint	<code>https://api.groq.com/openai/v1/chat/completions</code>	OpenAI-compatible chat completions API
Message Format	system + history + user	Full multi-turn conversation with system context

4.6 Rich Markdown Rendering

Bot responses are rendered with a custom markdown parser that converts: `###` headers to styled headings in gold, `**bold**` to gold-colored strong text, and newlines to `<br>` tags. All output is XSS-safe with HTML entity encoding.

5. Chatbot Consultation Handoff & Guardrails

5.1 Keyword-Based Intent Detection

After each LLM response, the backend scans both the user's message and the AI's reply for consultation-trigger keywords:

TRIGGER KEYWORD	SCANNED IN	ACTION
kundali	User message + AI reply	Sets <code>consultation_handoff: true</code>
birth chart	User message	Sets <code>consultation_handoff: true</code>
astrologer	User message + AI reply	Sets <code>consultation_handoff: true</code>
horoscope	User message	Sets <code>consultation_handoff: true</code>
dasha	User message	Sets <code>consultation_handoff: true</code>
matchmaking	User message	Sets <code>consultation_handoff: true</code>

5.2 Dual Handoff UI Cards

When `consultation_handoff` is true, the frontend renders two escalation cards:

- Primary Card (Dark Gradient):** Full-width card with Video icon, "Personal Kundali & Chart Reading Needed?" header, description of 1-on-1 Vedic Astrologer consultation, and "Book Live Video Consultation" CTA button with amber gradient styling.
- Secondary Card (Light):** Soft gradient card with UserCheck icon, "Astrologer Consultation Recommended" header, guidance text, and "Book Live Consultation" CTA button that navigates to `/consult`.

5.3 Response Structure

Every chatbot API response includes the following JSON structure:

```
{
  "reply": "Namaste! Here is your sacred guidance...",
  "recommendations_injected": true,
  "products": [
    {
      "id": 42,
      "slug": "shree-yantra-gold",
      "name": "Shree Yantra (Gold Plated)",
      "img": "https://lh3.googleusercontent.com/d/...",
      "desc": "Temple-energized sacred geometry tool",
      "price": "₹2,499.00",
      "raw_price": 249900
    }
  ],
  "consultation_handoff": false
}
```

5.4 Content Guardrails

- **Medical/Legal Advice:** AI declines and reframes as spiritual guidance.
 - **Deterministic Claims:** Phrases like "guarantees wealth" or "cures all problems" are strictly forbidden in the system prompt.
 - **Astrological Claims:** Must be phrased conversationally without outcome guarantees (e.g., "supports positive cosmic alignment" not "will make you rich").
-

6. Hybrid Recommendation Engine Architecture

The recommendation pipeline combines Gorse ML collaborative filtering with custom Astrological Soft-Boosting and visual Explicit Reasoning Badges. This hybrid approach ensures that recommendations are both behaviorally relevant (ML) and spiritually aligned (Vedic astrology).

6.1 End-to-End Pipeline

HYBRID SCORING PIPELINE

1. Candidate Retrieval (Gorse /api/recommend OR /api/popular)

2. Product Hydration (Supabase PostgreSQL full product details)

3. Hard Filter: Stock check (in_stock = true, stock > 0)

4. Soft Boost 1: +0.20 for matching Kundali Mahadasha Category

5. Soft Boost 2: +0.15 for matching Astrological Remedy Keyword

6. Reasoning Badge Attachment (dynamic based on score thresholds)

7. Sort by composite astroScore descending

8. Output Top Candidates to Shop Page, PDP Carousel & Chat

6.2 Two Recommendation Endpoints

ENDPOINT	METHOD	PURPOSE	GORSE QUERY
/api/recommendations/:userId	GET	User-to-Item personalized recs	GET /api/recommend/{userId}?n=4
/api/recommendations/item/:itemId	GET	Item-to-Item similarity (PDP)	GET /api/item/{itemId}/neighbors?n=4

6.3 User-to-Item Flow (GET /:userId)

1. Query Gorse for personalized recommendations (/api/recommend/{userId}?n=4).

2. If Gorse returns empty (new user), fallback to Gorse popular items (/api/popular?n=4).

3. Hydrate product IDs from Supabase PostgreSQL with full product details.

4. If still empty (cold start), fallback to Supabase top-selling bestsellers sorted by reviews and rating.

5. Check global Kundali profile cache for the user.

6. If profile exists, apply scoreAndRankProducts() for astrological soft-boosting.

7. Return final ranked products with reasonBadge and astro_boosted flag.

6.4 Item-to-Item Flow (GET /item/:itemId)

1. Query Gorse for item neighbors (/api/item/{itemId}/neighbors?n=4).

2. Hydrate product IDs from Supabase.

3. Attach default reasonBadge: "Frequently Bought Together" to all results.

7. Gorse ML Engine & Collaborative Filtering

Gorse is an open-source recommendation engine running as a Docker container (`zhenghaoz/gorse-in-one`) that provides real-time collaborative filtering based on user behavioral signals.

7.1 Docker Deployment

PARAMETER	VALUE
Docker Image	<code>zhenghaoz/gorse-in-one</code>
API Port	8088 (HTTP REST API)
Dashboard Port	8088/overview (Web UI)
Internal Cache	Redis (Port 6379, embedded)
Environment Variable	<code>GORSE_URL=http://localhost:8088</code>

7.2 Gorse API Endpoints Used

GORSE ENDPOINT	USED BY	PURPOSE
<code>GET /api/recommend/{userId}?n=N</code>	<code>recommendations.js</code> , <code>chat.js</code>	Personalized item recommendations
<code>GET /api/popular?n=N</code>	<code>recommendations.js</code>	Trending/popular items fallback
<code>GET /api/item/{itemId}/neighbors?n=N</code>	<code>recommendations.js</code>	Item-to-item similarity
<code>POST /api/feedback</code>	<code>telemetry.js</code>	Submit user behavioral events

7.3 Multi-Signal Learning

The ML engine learns from three weighted behavioral signals:

- `view_product` **(Implicit)**: Tracked when a user clicks a product card or visits a Product Detail Page. Lowest signal weight.
- `add_to_cart` **(Semi-Explicit)**: Tracked when a user adds an item to their shopping cart. Medium signal weight.
- `purchase` **(Explicit)**: Tracked when a transaction is completed. Highest signal weight.

Each signal is submitted to Gorse as a timestamped feedback entry: `{ FeedbackType, UserId, ItemId, Timestamp }`.

7.4 ID Format Handling

Gorse stores item IDs as strings, while Supabase uses integer IDs. The recommendation route handles this mismatch by converting all Gorse IDs through `Number(id)` casting and filtering out `NaN` values before querying Supabase with the `.in("id", numericIds)` filter.

When a brand-new visitor lands on Aroham for the first time, Gorse ML has zero click history for them. Instead of showing an empty area or breaking the layout, the backend runs a 3-tier fallback that guarantees products are always displayed.

Tier 1 (Personalized): Query Gorse `/api/recommend/{userId}?n=4` . If items are returned, use them.

Tier 3 (Database Bestsellers): If both Gorse queries fail or return empty, query Supabase directly: `SELECT * FROM products ORDER BY reviews DESC, rating DESC LIMIT 4.`

FALLBACK TIER	SOURCE	DEFAULT REASON	BADGE
Tier 1	Gorse Personalized	🌟 Curated For You	
Tier 2	Gorse Popular	🌟 Curated For You	
Tier 3	Supabase Bestsellers	🔥 Top Selling Bestseller	

The "Specially Curated For You" banner on the Shop Page is always populated with at least 4 product cards regardless of user history. This ensures the recommendation section never shows an empty state or loading skeleton indefinitely.

After the 3-tier fallback resolves, the system checks if the user has a cached Kundali profile. If they do, `scoreAndRankProducts()` is applied to re-rank even the fallback products by astrological relevance. This means even bestseller fallback products get re-ordered by Dasha alignment.

9. Astrological Soft-Boosting & Kundli Scoring

The `services/kundliScoring.js` module implements the Vedic astrological layer that bridges Kundli PDF generation with the product recommendation pipeline.

9.1 Derived Feature Extraction (`getDerivedAstrologyProfile()`)

When a user generates their Kundli PDF, the backend extracts non-PII derived features from the birth chart data:

- **mahadasha:** The user's active planetary period (e.g., Saturn, Rahu, Jupiter).
- **moonSign:** The zodiac sign of the Moon at birth (e.g., Aries, Taurus).
- **nakshatra:** The lunar mansion of the Moon at birth (e.g., Ashwini, Rohini).
- **recommendedCategories:** Array of product categories aligned with the active Dasha.
- **recommendedKeywords:** Array of product name/description keywords for matching.
- **guidanceNote:** A human-readable guidance text for the LLM system prompt.

9.2 Global Profile Cache

Derived profiles are stored in `global.kundaliProfiles` (in-memory) keyed by both user name and userId. This allows the recommendation engine and chatbot to access astrological context without re-querying the Kundli API. The cache persists for the lifetime of the Node.js process.

9.3 Scoring Algorithm (`scoreAndRankProducts()`)

Each product candidate starts with a base score of `1.0`. The algorithm applies the following transformations:

CHECK	CONDITION	SCORE IMPACT
Hard Filter	<code>in_stock === false</code> OR <code>stock ≤ 0</code>	Score set to <code>0</code> (eliminated)
Soft Boost 1	Product category matches any <code>recommendedCategories</code>	<code>+0.20</code>
Soft Boost 2	Product name/desc contains any <code>recommendedKeywords</code>	<code>+0.15</code>

9.4 Score Thresholds & Badge Assignment

SCORE RANGE	MEANING	REASONBADGE
<code>> 1.15</code>	Both category + keyword match	Aligned with {Dasha} Dasha
<code>> 1.0</code> and <code>≤ 1.15</code>	Category match only	Recommended for {Rashi} Rashi
<code>= 1.0</code>	No astrological match	Popular Choice
<code>= 0</code>	Out of stock	Filtered out entirely

After scoring, products are sorted by `astroScore` descending. Out-of-stock items (score = 0) are filtered out entirely and never surface in any recommendation context.

10. Complete Mahadasha-to-Category Mapping Reference

The `dashaMap` in `kundliScoring.js` maps each of the 9 Vimshottari Mahadasha planetary periods to specific product categories, remedy keywords, and guidance notes:

MAHADASHA	RECOMMENDED CATEGORIES	REMEDY KEYWORDS
Saturn	Yantras, Rudraksha, Protection	Shani, Saturn, 8-Mukhi, Neelam, Protection, Iron, Shree Yantra
Rahu	Protection, Yantras, Rudraksha	Rahu, 9-Mukhi, Hessonite, Gomed, Protection, Kavach
Ketu	Rudraksha, Protection, Puja & Yantras	Ketu, Cat's Eye, Lehsuniya, 9-Mukhi, Moksha
Jupiter	Gemstones, Yantras, Rudraksha	Guru, Jupiter, Yellow Sapphire, Pukhraj, 5-Mukhi, Kanakdhara
Venus	Jewelry & Malas, Gemstones, Yantras	Shukra, Venus, Diamond, Opal, Sphatik, Lakshmi
Mars	Yantras, Protection, Rudraksha	Mangal, Mars, Red Coral, Moonga, 3-Mukhi, Hanuman
Sun	Yantras, Rudraksha, Gemstones	Surya, Sun, 1-Mukhi, Ruby, Manik, Aditya
Moon	Jewelry & Malas, Gemstones, Rudraksha	Chandra, Moon, Pearl, Moti, 2-Mukhi, Moonstone, Silver
Mercury	Gemstones, Yantras, Rudraksha	Budh, Mercury, Emerald, Panna, 4-Mukhi, Ganesh

10.1 Guidance Notes per Dasha

- Saturn:** "Saturn Mahadasha requires discipline, patience, and protective grounding remedies."
- Rahu:** "Rahu Mahadasha requires clarity, protection from illusions, and stabilizing Yantras."
- Ketu:** "Ketu Mahadasha favors spiritual inward reflection and protective Kavach remedies."
- Jupiter:** "Jupiter Mahadasha brings wisdom, expansion, and prosperity when supported by blessed Yantras."
- Venus:** "Venus Mahadasha invites harmony, beauty, and abundance through energized Sphatik & Lakshmi remedies."
- Mars:** "Mars Mahadasha calls for courage, vitality, and balancing intense cosmic energy."
- Sun:** "Sun Mahadasha enhances leadership, vitality, and solar alignment through Surya remedies."
- Moon:** "Moon Mahadasha enhances emotional peace, intuition, and mental tranquility."
- Mercury:** "Mercury Mahadasha sharpens intellect, communication, and commercial success."

11. Explicit Reasoning Badges System (Point 9 Deep Dive)

To eliminate recommendation ambiguity and increase buyer trust, the system attaches human-readable visual badges directly to product cards based on scoring origin. Users can see *why* a product is recommended to them.

11.1 Visual Badge Taxonomy

- 🌟 `Aligned with [Dasha] Dasha` — Attached to items soft-boosted by active Mahadasha remedies. Highest astrological relevance (score > 1.15). Example: "🌟 Aligned with Saturn Dasha".
- 🌙 `Recommended for [Rashi] Rashi` — Attached to items matching Moon sign remedies. Medium astrological relevance (score > 1.0). Example: "🌙 Recommended for Aries Rashi".
- 💎 `Frequently Bought Together` — Attached to item-to-item neighbor recommendations on Product Detail Pages. Indicates behavioral co-purchase patterns.
- 🔥 `Popular Choice` — Attached to Gorse ML recommendations that don't match astrological profile. Indicates community trending.
- 🔥 `Top Selling Bestseller` — Attached to Supabase fallback bestsellers. Indicates overall sales performance.
- 🌟 `Curated For You` — Default badge for personalized and popular Gorse recommendations before astrological scoring.

11.2 Badge Assignment Logic

Badges are assigned at three different points in the pipeline:

- Pre-Scoring Default:** All products from Gorse get 🌟 `Curated For You`. Bestseller fallbacks get 🔥 `Top Selling Bestseller`. Item neighbors get 💎 `Frequently Bought Together`.
- Post-Scoring Override:** If astrological scoring is applied via `scoreAndRankProducts()`, the default badge is overridden based on the final score threshold.
- Frontend Rendering:** The badge is rendered as a pill-shaped tag at the top of each product card across all surfaces.

11.3 Rendering Surfaces

SURFACE	COMPONENT	BADGE STYLE
Shop Page Banner	<code>ShopPage.tsx</code>	Pill with amber-200 text on amber-500/20 background
Product Detail Page	<code>ProductDetailPage.tsx</code>	Pill with amber-800 text on amber-100 background
Chatbot Product Cards	<code>AstroChatWidget.tsx</code>	Inline with product mini-cards

12. Recommendation Surface Points (Shop, PDP, Chat)

Recommendations are surfaced across three distinct touchpoints in the user journey, each with specialized UI treatment:

12.1 Shop Page — "Specially Curated For You" Banner

- **Trigger:** Loaded on Shop Page mount via `GET /api/recommendations/{userId}`.
- **Visibility:** Only shown when no active category/price filters are applied.
- **Layout:** Full-width gradient banner (maroon-to-dark) with 2×4 product grid.
- **Product Cards:** Glass-morphism cards with backdrop blur, product image (hover zoom), name, short description, price in amber, reason badge pill, and "View" button.
- **Header:** " ✨ Specially Curated For You" with sparkle icon and subtitle.
- **Responsive:** 2 columns on mobile, 4 columns on desktop.

12.2 Product Detail Page — "Frequently Bought Together" Carousel

- **Trigger:** Loaded on PDP mount via `GET /api/recommendations/item/{itemId}`.
- **Layout:** 2×4 grid below the main product details section.
- **Product Cards:** White rounded cards with shadow, product image (hover scale), name, short description, price, and reason badge.
- **Header:** "Frequently Bought Together" with Sparkles icon and "Devotees who brought home this sacred item also resonated with these."
- **Badge:** All items get `💎 Frequently Bought Together` by default.

12.3 AstroGuide Chat — Inline Product Mini-Cards

- **Trigger:** Loaded as part of every chat response via `POST /api/chat`.
- **Layout:** Stacked vertical list below the bot's text response.
- **Product Cards:** Compact cards with 44×44px thumbnail, product name (truncated), price in amber, "View Product" button (ExternalLink icon), and "Add to Cart" button (ShoppingBag icon).
- **Add All Button:** "Add All Recommended to Cart (Save 10%)" — appears when 2+ products are shown.
- **Section Header:** " ✨ Recommended Sacred Remedies" in tracking-widest uppercase.

13. Kundli PDF Generation Engine (18+ D-Charts)

Users can click " 🏠 Make My Kundli " in the top navigation header to generate and download a comprehensive 16+ page Vedic Kundli PDF report.

13.1 Technical Architecture

- **Endpoint:** `POST /api/kundli/generate` .
- **Request Payload:** `{ name, gender, location, date (DD/MM/YYYY), time (HH:MM), userId }` .
- **Astronomical Calculations:** VedAstro API (`api.vedastro.org`) for exact geographic coordinates, 4-decimal lat/long, GMT timezone, and 12-planet/house positions.
- **18+ North Indian Divisional Charts:** SVG/HTML rendering for D1, D2, D3, D4, D7, D9, D10, D12, D16, D20, D24, D27, D30, D40, D45, D60, Surya Chart, Chandra Chart, and KP Chalit.
- **Precision Astronomical Tables:** Ascendant Lagna top row, Panchang, Avakhada Chakra, KP Cusps, SAV points, BAV tables, Shadbala 6-fold strength, Bhavbala house ranks, and Vimshottari/Yogini dasha timelines with visual `CURRENT` status badges.
- **Instant Browser Download:** Compiled via Puppeteer and EJS template (`kundli_report.ejs`), streamed with `Content-Type: application/pdf` .

13.2 6-Stage Processing Pipeline

1. **Geographic Resolution:** Location string is sent to VedAstro `AddressToGeoLocation` API. Returns exact latitude (4 decimals), longitude, and place name. Falls back to Delhi coordinates (28.6139, 77.2090) if API fails.
2. **Dynamic Timezone Calculation:** Uses `geo-tz` library to compute exact timezone offset from lat/long. Combined with `moment-timezone` for DST-aware offset formatting (e.g., "+05:30").
3. **Parallel VedAstro Queries (11 concurrent):** All planet data, all house data, Tithi, Nithya Yoga, Karana, Sunrise, Sunset, Moon Nakshatra, and outer planets (Uranus, Neptune, Pluto) are fetched simultaneously via `Promise.all()` .
4. **Avakhada Chakra Computation:** Moon Nakshatra, Charan, Name Alphabet (syllable lookup), Sign, Sign Lord, Paya (Loha/Tamra/Rajat/Swarna based on moon house), Yunja, Gan, Nadi, Yoni, Varna, and Tatva are computed using the `avakhada_lookup.json` reference table.
5. **Derived Profile Caching:** The active Mahadasha and Moon Sign are extracted and cached globally via `getDerivedAstrologyProfile()` for recommendation boosting.
6. **PDF Compilation:** All chart data is passed to `compileAstrologyPdf()` which renders the EJS template with Puppeteer and returns a PDF buffer streamed directly to the browser.

14. Kundli Astronomical Calculations & Avakhada Chakra

14.1 VedAstro API Endpoints (11 Parallel Queries)

QUERY	VEDASTRO ENDPOINT	RETURNS
All Planets	/AllPlanetData/PlanetName/All/ ...	12 planet positions, signs, degrees, lords
All Houses	/AllHouseData/HouseName/All/ ...	12 house cusps, signs, lords
Tithi	/Tithi/ ...	Lunar day in Panchang
Nithya Yoga	/NithyaYoga/ ...	Yoga of the day
Karana	/Karana/ ...	Half-lunar day
Sunrise	/SunriseTime/ ...	Local sunrise time
Sunset	/SunsetTime/ ...	Local sunset time
Moon Nakshatra	/MoonNakshatra/ ...	Moon's asterism and pada
Uranus	/AllPlanetData/PlanetName/Uranus/ ...	Outer planet position
Neptune	/AllPlanetData/PlanetName/Neptune/ ...	Outer planet position
Pluto	/AllPlanetData/PlanetName/Pluto/ ...	Outer planet position

14.2 Avakhada Chakra Derived Fields

FIELD	DERIVATION SOURCE	EXAMPLE
Nakshatra	Moon Nakshatra API response	Rohini
Charan (Pada)	Parsed from Nakshatra string (after hyphen)	2
Name Alphabet	Syllable lookup from avakhada_lookup.json	Vi
Sign (Rashi)	Moon planet data PlanetRasiD1Sign.Name	Taurus
Sign Lord	Moon planet data PlanetLord0fZodiacSign.Name	Venus
Paya	Moon house position relative to Lagna	Swarna (Gold)
Yunja	Nakshatra lookup table	Madhya
Gan	Nakshatra lookup table	Manushya
Nadi	Nakshatra lookup table	Madhya
Yoni	Nakshatra lookup table	Sarpa

FIELD	DERIVATION SOURCE	EXAMPLE
Varna	Rashi lookup table	Vaishya
Tatva	Rashi lookup table	Prithvi

14.3 Paya Calculation Logic

Paya (foot) is determined by the Moon's house position relative to the Ascendant Lagna:

- **Loha (Iron):** Moon in Houses 1, 6, 11
- **Tamra (Copper):** Moon in Houses 3, 7, 10
- **Rajat (Silver):** Moon in Houses 2, 5, 9
- **Swarna (Gold):** Moon in Houses 4, 8, 12

14.4 Kundli-to-Recommendation Bridge

After PDF compilation, the `SignLord` is used as the Mahadasha identifier. If SignLord is "Saturn", the Mahadasha is set to Saturn. Otherwise, the SignLord value is used directly. This bridges the astrological calculation to the recommendation scoring pipeline.

15. Telemetry & Event Ingestion Pipeline

The telemetry system captures user behavioral signals and forwards them to the Gorse ML engine for real-time model training.

15.1 Telemetry Endpoints (`routes/telemetry.js`)

ENDPOINT	METHOD	PAYLOAD	GORSE FEEDBACKTYPE
<code>/api/telemetry/click</code>	POST	<code>{ userId, productId }</code>	<code>view_product</code>
<code>/api/telemetry/event</code>	POST	<code>{ userId, productId, eventType }</code>	Dynamic (<code>add_to_cart</code> , <code>purchase</code> , etc.)

15.2 Feedback Submission Format

Both endpoints construct a Gorse-compatible feedback array and POST it to `GORSE_URL/api/feedback` :

```
[{
  "FeedbackType": "view_product",
  "UserId": "user_demo_123",
  "ItemId": "42",
  "Timestamp": "2026-07-28T18:30:00.000Z"
}]
```

15.3 Frontend Trigger Points

EVENT	FRONTEND TRIGGER	BACKEND ROUTE	GORSE FEEDBACKTYPE
Product View	Clicking a product card / visiting PDP	<code>POST</code> <code>/api/telemetry/click</code>	<code>view_product</code>
Add to Cart	Clicking "Add to Cart" on PDP	<code>POST</code> <code>/api/telemetry/event</code>	<code>add_to_cart</code>
Purchase	Completing a transaction	<code>POST</code> <code>/api/telemetry/event</code>	<code>purchase</code>

15.4 Error Handling

Both telemetry endpoints implement graceful error handling. If the Gorse container is unreachable, the endpoint returns `{ success: false, error: "..." }` without throwing a 500 error. This ensures telemetry failures never break the user's shopping experience.

15.5 Debug Logging

The backend captures all `console.log` and `console.error` calls in a circular `global.debugLogs` buffer (max 200 entries) with timestamps. This enables real-time debugging without external log aggregation services.

16. Complete API Reference

16.1 Chat API

FIELD	DETAILS
Endpoint	POST /api/chat
Request Body	{ message: string, userId?: string, pageContext?: string, history?: Array }
Response	{ reply: string, recommendations_injected: boolean, products: Array, consultation_handoff: boolean }
LLM Provider	Groq API (llama-3.3-70b-versatile)
Error Response	{ error: string, details: string } (HTTP 500)

16.2 Recommendations API

FIELD	DETAILS
Endpoint 1	GET /api/recommendations/:userId
Response 1	{ success: boolean, recommendations: Array, astro_boosted: boolean }
Endpoint 2	GET /api/recommendations/item/:itemId
Response 2	{ recommendations: Array }

16.3 Kundli API

FIELD	DETAILS
Endpoint	POST /api/kundli/generate
Request Body	{ name: string, gender: string, location: string, date: string (DD/MM/YYYY), time: string (HH:MM), userId?: string }
Response	Binary PDF stream (Content-Type: application/pdf)
Headers	Content-Disposition: attachment; filename="Kundli-{name}.pdf"

16.4 Telemetry API

FIELD	DETAILS
Click Endpoint	POST /api/telemetry/click
Event Endpoint	POST /api/telemetry/event

FIELD	DETAILS
Request Body	<code>{ userId: string, productId: string, eventType?: string }</code>
Response	<code>{ success: boolean, gorseResponse: string }</code>

16.5 Health Check

FIELD	DETAILS
Endpoint	<code>GET /api/health</code>
Response	<code>{ status: "ok", service: "aroham-backend" }</code>

17. Integration Points & Tech Stack

Component	Technology	Details
Frontend	React 19 + Vite 6	TypeScript, TailwindCSS, Lucide icons
Backend	Express.js (Node.js)	Port 5000, CORS enabled, 50MB body limit
LLM Provider	Groq API	Model: llama-3.3-70b-versatile
ML Engine	Gorse (Docker)	zhenghaoz/gorse-in-one , Port 8088
Kundli API	VedAstro API	Geo-location & planetary calculations
PDF Renderer	Puppeteer + EJS	Multi-page A4 PDF compiler
Product DB	Supabase PostgreSQL	Cloud-hosted products, users, orders
Authentication	Custom Auth Module	User sessions and guest tracking
Payments	Razorpay	Webhook signature verification via rawBody
Shipping	Shippo API	Logistics & order tracking
Image Hosting	Google Drive + LH3	Auto-converted to direct URLs
Timezone	geo-tz + moment-timezone	Dynamic timezone from lat/long
State Management	Shared Cart State	Cross-component cart operations
Astro Scoring	kundliScoring.js	Custom Vedic soft-boost engine

17.1 Security Headers

The Express backend applies the following security headers to all responses:

- X-Content-Type-Options: nosniff — Prevents MIME type sniffing.
- X-Frame-Options: DENY — Prevents clickjacking via iframe embedding.
- Cache-Control: no-store, no-cache, must-revalidate, proxy-revalidate — Disables all caching for API responses.

17.2 Environment Variables

Variable	Default	Used By
PORT	5000	Express server
GROQ_API_KEY	(required)	Chat API
LLM_MODEL	llama-3.3-70b-versatile	Chat API
GORSE_URL	http://localhost:8088	Recommendations, Chat, Telemetry

VARIABLE	DEFAULT	USED BY
SUPABASE_URL	(required)	All DB operations
SUPABASE_KEY	(required)	All DB operations

18. System Boundaries & Edge Cases

SCENARIO	WHAT HAPPENS	USER EXPERIENCE
User asks real Kundali reading in chat	Handoff card appears with "Book Live Consultation"	Graceful escalation to human astrologer
User asks medical/legal advice	AI declines and reframes as spiritual guidance	Polite redirect within Vedic context
Groq API is down	Error message: "I had trouble reading the stars"	Chat remains functional, no AI reply
Gorse container is down	Fallback to Supabase top-sellers	Recommendations still appear, just not personalized
Brand new user (cold start)	3-tier fallback to Supabase bestsellers	Always shows products, never empty
User generates Kundli	18+ D-chart PDF downloads + Dasha profile cached	Instant PDF download + recommendation boost
VedAstro API is unreachable	Falls back to Delhi coordinates (28.6139, 77.2090)	PDF still generates with approximate data
Telemetry endpoint fails	Returns { success: false } silently	Shopping experience unaffected
User has no Kundli profile	Astrological boosting is skipped	Standard Gorse ML recommendations shown
Out-of-stock product in Gorse results	Hard-filtered with score = 0	Never shown to users

19. Final System Readiness Sign-Off

Comprehensive system verification confirmed 100% operational readiness across all microservices:

- **AI Chatbot System:** 100% Operational (Groq LLaMA 3.3 70B + Multi-Turn Memory + System Prompt Engineering + Product Injection + Handoff Card + Add All to Cart Bundle + Quick Suggestion Pills + Proactive Bubble + Guest Tracking).
- **Recommendation Engine:** 100% Operational (Gorse ML Collaborative Filtering + 3-Tier Cold Start Fallback + Astrological Soft Boost + Explicit Reason Badges + User-to-Item + Item-to-Item + 3 Surface Points).
- **Kundli PDF Generator:** 100% Operational (18+ D-Charts + 11 Parallel VedAstro Queries + Avakhada Chakra + Dynamic Timezone + Kundli-to-Recommendation Bridge + Instant PDF Download).
- **Telemetry Pipeline:** 100% Operational (3 Signal Types + Gorse Feedback Ingestion + Graceful Error Handling + Debug Log Buffer).
- **Overall System Readiness: Production Ready.**